Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 6 DISC...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 6
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

A company uses Amazon SageMaker for its ML pipeline in a production environment. The company has large input data sizes up to 1 GB and processing times up to 1 hour. The company needs near real-time latency.
Which SageMaker inference option meets these requirements?

A. Real-time inference

B. Serverless inference

C. Asynchronous inference

D. Batch transform

Show Suggested Answer

---

by galliaj at *Nov. 1, 2024, 7:10 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

jove Highly Voted 1 year, 4 months ago

Selected Answer: C

Real-Time Inference: Immediate responses for high-traffic, low-latency applications.
>> Asynchronous Inference: Near real-time for large payloads and longer processing.
Batch Transform: Large-scale, offline processing without real-time needs.
Serverless Inference: Low-latency inference for intermittent or unpredictable traffic without managing infrastructure.

upvoted 19 times

Salmariaz Most Recent 2 months, 4 weeks ago

Selected Answer: C

A) Real-Time (not correct)– Immediate response & best for chatbots, fraud detection, live recommendations.
B) Serverless Inference (not correct)– Real-time without managing servers & scales automatically for low/unpredictable traffic.
C) Asynchronous (correct) – Queued requests & best for heavy or variable-latency workloads(almost real time).
D) Batch Transform – (not correct) Processes large files & results stored for large-scale predictions.
Multi-Model Endpoints – Serve multiple models via one endpoint for multi-tenant apps.

upvoted 1 times

d8945a1 5 months, 3 weeks ago

Selected Answer: C

The company needs NEAR real-time, not true real-time.

upvoted 1 times

virtual_vikram 6 months, 1 week ago

Answer : C
Real time inference : Doesn't support input size more than 6 MB.
Async : Designed for near real-time for larger input size.

upvoted 1 times

muimijustus 6 months, 2 weeks ago

Selected Answer: C

Amazon SageMaker Asynchronous Inference is a capability in SageMaker AI that queues incoming requests and processes them asynchronously. This option is ideal for requests with large payload sizes (up to 1GB), long processing times (up to one hour), and near real-time latency requirements.

upvoted 1 times

Rayenwalid 7 months ago

Selected Answer: C

Asynchronous inference is a nice balance between speed and efficiency for this use case.

upvoted 1 times

65703c1 8 months, 1 week ago

Selected Answer: C

C is the correct answer.

upvoted 1 times

Dimkaaa 8 months, 3 weeks ago

Selected Answer: C

Amazon SageMaker Asynchronous Inference is a capability in SageMaker AI that queues incoming requests and processes them asynchronously. This option is ideal for requests with large payload sizes (up to 1GB), long processing times (up to one hour), and near real-time latency requirements.

https://docs.aws.amazon.com/sagemaker/latest/dg/async-inference.html

upvoted 2 times

vm74 11 months ago

Selected Answer: C

Real-Time Inference
Latency Low, immediate
Payload Size Typically smaller
Processing Time Typically short

Asynchronous Inference
Latency Delayed, can be in minutes
Payload Size Can be large, up to 1GB
Processing Time Can be long, up to 15 minutes

Since the requirement is near realtime upto 1 hour, and payload size upto 1 GB, Option C is appropriate selection

upvoted 1 times

INDKAR 11 months, 3 weeks ago

Selected Answer: A

Asynchronous inference can process data size of 1 GB

upvoted 2 times

tcl08 12 months ago

Selected Answer: A

Asynchronous inference processes requests in the background, returning a response ID and allowing the client to check for results later, while real-time inference delivers redictions with minimal delay, suitable for interactive applications

upvoted 1 times

Rcosmos 12 months ago

Selected Answer: C

Explicação:A inferência assíncrona do Amazon SageMaker é ideal quando:
Os dados de entrada são grandes (por exemplo, até 1 GB),O tempo de processamento Pode ser longo (até 1 hora por solicitação),E você ainda precisa de respostas com latência razoável, mas não exige resposta instantânea como em APIs síncronas.
Ela permite que você envie a solicitação, continue processando outras tarefas e recupere o resultado quando estiver pronto, o que evita timeouts comuns em inferência síncrona.

upvoted 1 times

Amar949499 1 year, 1 month ago

Selected Answer: C

Here the keyword is near “real time latency “
Asynchronous inference. queues incoming requests and processes them asynchronously. This option is ideal for requests with large payload sizes (up to 1GB), long processing times (up toAsynchronous Inference one hour), and near real-time latency requirements

upvoted 1 times

Nopnov 1 year, 1 month ago

Selected Answer: C

Amazon SageMaker Asynchronous Inference is a capability in SageMaker AI that queues incoming requests and processes them asynchronously. This option is ideal for requests with large payload sizes (up to 1GB), long processing times (up to one hour), and near real-time latency requirements

upvoted 1 times

JJwin 1 year, 1 month ago

Selected Answer: A

Real-time inference in Amazon SageMaker is designed for low-latency, high-throughput applications where predictions need to be made immediately after data is processed. Since the company requires near real-time latency for their ML pipeline and has processing times of up to 1 hour and input sizes up to 1 GB, real-time inference is the most suitable option.
With real-time inference, you can deploy your trained models as an API endpoint and get predictions on demand, ensuring low latency. This is ideal for situations where you need immediate responses after submitting the data.

upvoted 2 times

Willdoit 1 year, 1 month ago

Selected Answer: A

The company requires near real-time latency, which means the model needs to respond quickly to inference requests. Real-time inference in Amazon SageMaker is designed for low-latency applications where predictions are needed in milliseconds to seconds.
C. Asynchronous inference – Useful for large requests that take minutes or hours to process, but it is not real-time.

upvoted 2 times

Jessiii 1 year, 1 month ago

Selected Answer: C

Asynchronous inference in Amazon SageMaker is ideal when you have large input data sizes (like the 1 GB mentioned) and relatively long processing times (like up to 1 hour). While real-time inference typically offers lower latency, it may struggle with large datasets or complex models that require more processing time. In contrast, asynchronous inference can handle large inputs and longer processing times without needing immediate results. It processes the data and provides the results later, which might be acceptable if your requirement for near real-time latency can be slightly relaxed (for instance, if results can be retrieved within minutes rather than immediately).

upvoted 2 times