## Exam AWS Certified AI Practitioner AIF-C01 All Questions

**View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam**

Go to Exam


## EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 10 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 10
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

---

A company wants to use language models to create an application for inference on edge devices. The inference must have the lowest latency possible.
Which solution will meet these requirements?

A. Deploy optimized small language models (SLMs) on edge devices.

B. Deploy optimized large language models (LLMs) on edge devices.

C. Incorporate a centralized small language model (SLM) API for asynchronous communication with edge devices.

D. Incorporate a centralized large language model (LLM) API for asynchronous communication with edge devices.

Show Suggested Answer

---

by galliaj at *Nov. 1, 2024, 7:19 p.m.*


Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.


## Comments

**Moon** Highly Voted 1 year, 3 months ago

Selected Answer: A

A: Deploy optimized small language models (SLMs) on edge devices.

Explanation:
Deploying optimized small language models (SLMs) on edge devices ensures low latency because the inference happens directly on the device without relying on cloud communication. Small language models are lightweight and designed to run efficiently on devices with limited resources, making them ideal for edge computing.

upvoted 8 times

**65703c1** Most Recent 8 months, 1 week ago

Selected Answer: A

A is the correct answer.

upvoted 1 times

**PrahladB** 9 months, 3 weeks ago

Selected Answer: A

Verified its A. Use AI for solution.

upvoted 1 times

**Nagen_007** 10 months, 2 weeks ago

Selected Answer: A

Deploy optimized small language models (SLMs) on edge devices.

upvoted 1 times

**INDKAR** 11 months, 3 weeks ago

Selected Answer: A

Lowest latency = Model has to be deployed on the Edge device
Only SLM can be deployed on Edge device due to capacity constraints

upvoted 3 times

**Rcosmos** 12 months ago

Selected Answer: U

Quando o objetivo é inferência com a menor latência possível, a melhor abordagem é executar o modelo diretamente no dispositivo de borda (edge).
SLMs (Small Language Models) são projetados para serem leves, rápidos e eficientes, o que os torna ideais para: Dispositivos com recursos limitados
Tempo de resposta imediato . Execução offline ou com pouca conectividade

upvoted 1 times

**Jessiii** 1 year, 1 month ago

Selected Answer: A

Optimized small language models (SLMs) are specifically designed to run efficiently on edge devices with limited resources (such as memory and processing power). Deploying smaller, optimized models directly on the edge devices allows for near-instantaneous inference with minimal latency, as the data doesn't need to travel to a central server for processing.

upvoted 1 times

**Aryan_10** 1 year, 3 months ago

Selected Answer: A

Lowest latency possible - SLM

upvoted 1 times

**Nicocacik** 1 year, 4 months ago

Selected Answer: A

Low latency with edge devices -> SLM

upvoted 1 times

**Blair77** 1 year, 4 months ago

A is good - Minimal latency: SLMs are designed to run efficiently on resource-constrained devices, offering fast inference directly on the device.

upvoted 1 times

**jove** 1 year, 4 months ago

Selected Answer: A

SLM on edge devices

upvoted 2 times

**tccusa** 1 year, 5 months ago

Selected Answer: A

SLM on edge devices is the correct solution.

upvoted 2 times

**galliaj** 1 year, 5 months ago

Using Optimized Small Language Models (SLMs) on edge devices is the best choice because they are designed to run efficiently within the resource constraints of edge hardware. This minimizes latency and helps deliver fast inference times while using less computational power and memory. The problem with trying to use centralized APIs is the associated latentcy.

upvoted 4 times