Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 66 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 66
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

A company wants to develop a large language model (LLM) application by using Amazon Bedrock and customer data that is uploaded to Amazon S3. The company's security policy states that each team can access data for only the team's own customers.
Which solution will meet these requirements?

A. Create an Amazon Bedrock custom service role for each team that has access to only the team's customer data.

B. Create a custom service role that has Amazon S3 access. Ask teams to specify the customer name on each Amazon Bedrock request.

C. Redact personal data in Amazon S3. Update the S3 bucket policy to allow team access to customer data.

D. Create one Amazon Bedrock role that has full Amazon S3 access. Create IAM roles for each team that have access to only each team's customer folders.

Show Suggested Answer

---

by jove at *Nov. 10, 2024, 3:31 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

459470f 7 months, 1 week ago

Selected Answer: A

I would say C is incorrect because it doesn't follow least privilege security. The single Bedrock role is a weak point and provides more access than is necessary

upvoted 1 times

65703c1 8 months, 1 week ago

Selected Answer: A

Ais the correct answer

upvoted 1 times

penguins2 9 months ago

Selected Answer: A

The trick of this question is Bedrock will use which role to access S3 data. It uses service role. For D, even each customer is restricted to a specific S3 bucket, once the customer uses Bedrock, Bedrock still uses all S3 data to generate the result.

upvoted 2 times

Rcosmos 10 months, 2 weeks ago

Selected Answer: D

resposta correta é D. Crie uma função do Amazon Bedrock que tenha acesso total ao Amazon S3. Crie funções do IAM para cada equipe que tenha acesso apenas às pastas de clientes de cada equipe.

Essa abordagem garante que o modelo de linguagem possa acessar todos os dados armazenados no Amazon S3, enquanto o controle de acesso é gerenciado por funções do IAM. Cada equipe terá acesso exclusivo apenas às pastas de clientes de sua responsabilidade, cumprindo a política de segurança da empresa.

upvoted 1 times

HarshulSinghThakur 1 year, 1 month ago

This option ensures that each team has a dedicated service role with permissions specifically tailored to access only their own customer data. This approach aligns with the company's security policy by enforcing strict access controls at the role level, ensuring that teams cannot access data belonging to other teams. By creating separate roles for each team, the solution adheres to the principle of least privilege, which is a fundamental security best practice. This method also simplifies auditing and management of access permissions, as each role's permissions can be reviewed and adjusted independently.

upvoted 1 times

Jessiii 1 year, 1 month ago

Selected Answer: A

A. Create an Amazon Bedrock custom service role for each team that has access to only the team's customer data.: This would require multiple service roles for Amazon Bedrock itself, which could lead to unnecessary complexity and overhead in role management. It's better to use IAM roles to control team-specific data access.

upvoted 4 times

pavankvv 1 year, 1 month ago

Selected Answer: D

By creating a single Bedrock role with full S3 access and then using IAM roles to control access to the customer data folders, the company can meet its requirements for developing the LLM application while also adhering to its security policy.

upvoted 2 times

Moon 1 year, 3 months ago

Selected Answer: A

A: Create an Amazon Bedrock custom service role for each team that has access to only the team's customer data.

Explanation:
To comply with the company's security policy requiring each team to access only their own customer data, the best approach is to create custom service roles in Amazon Bedrock for each team. These roles should have fine-grained permissions, granting access only to the specific Amazon S3 data (e.g., folders or buckets) associated with each team's customers. This ensures compliance with the principle of least privilege.

Wrong: D. Create one Amazon Bedrock role that has full Amazon S3 access. Create IAM roles for each team that have access to only each team's customer folders: Giving Bedrock full S3 access is a major security risk. Even with team-specific IAM roles, the Bedrock role could be exploited to access any data in S3.

upvoted 3 times

may2021_r 1 year, 3 months ago

Selected Answer: A

The correct answer is A. Custom service roles for each team provide granular control over customer data access.

upvoted 1 times

may2021_r 1 year, 3 months ago

Selected Answer: A

A. Create an Amazon Bedrock custom service role for each team that has access to only the team's customer data.

upvoted 1 times

Contactfornitish 1 year, 4 months ago

Selected Answer: D

A. Create an Amazon Bedrock custom service role for each team that has access to only the team's customer data
While this restricts data access, managing multiple service roles for Amazon Bedrock per team is unnecessarily complex and does not align with Bedrock's design of using a single service role.

B. Create a custom service role that has Amazon S3 access. Ask teams to specify the customer name on each Amazon Bedrock request
Relying on teams to specify the customer name without enforcing access control policies does not guarantee compliance with the security policy.

C. Redact personal data in Amazon S3. Update the S3 bucket policy to allow team access to customer data
Redacting personal data is helpful for privacy but does not solve the issue of restricting access based on team-specific customer data.

upvoted 3 times

fed6485 1 year, 4 months ago

Selected Answer: A

it has to be A.
B, absurd
C, would let all teams access all data, even if scrubbed/redacted ..
D, it would not solve the problem as Bedrock would have access, know and reply with the full knowledge of all customers, IAM roles for each team won't stop Bedrock from knowing and replying with that data..

A.
You can also create a custom service role and customize the attached permissions to your specific use-case. If you use the console, you can select this role instead of letting Amazon Bedrock create one for you.

https://docs.aws.amazon.com/bedrock/latest/userguide/security-iam-sr.html

upvoted 1 times

Blair77 1 year, 4 months ago

Selected Answer: A

Adherence to Principle of Least Privilege: By creating a custom service role for each team that grants access only to their specific customer data in S3, you ensure compliance with the principle of least privilege. Each team will have the minimum necessary permissions to access only their relevant data. D is wrong.

upvoted 1 times

taka5094 1 year, 4 months ago

Selected Answer: A

Creating a Bedrock role with access to all S3 data violates the principle of least privilege.

upvoted 2 times

jove 1 year, 4 months ago

I think it should be D, one IAM role for the service, and multiple IAM roles for the teams

upvoted 3 times

urbanmonk 1 year, 3 months ago

D makes sense on the surface but it talks about distinct customer's folders 📂 , which was not mentioned in the question. And granting Bedrock Full S3 access is certainly a huge red flag. So the answer cannot be D. That leaves "A" as the only plausible solution and answer.

upvoted 1 times