Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 67 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 67
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

A medical company deployed a disease detection model on Amazon Bedrock. To comply with privacy policies, the company wants to prevent the model from including personal patient information in its responses. The company also wants to receive notification when policy violations occur.
Which solution meets these requirements?

A. Use Amazon Macie to scan the model's output for sensitive data and set up alerts for potential violations.

B. Configure AWS CloudTrail to monitor the model's responses and create alerts for any detected personal information.

C. Use Guardrails for Amazon Bedrock to filter content. Set up Amazon CloudWatch alarms for notification of policy violations.

D. Implement Amazon SageMaker Model Monitor to detect data drift and receive alerts when model quality degrades.

Show Suggested Answer

---

by jove at *Nov. 10, 2024, 3:41 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

Jessiii Highly Voted 1 year, 1 month ago

Selected Answer: C

C. Use Guardrails for Amazon Bedrock to filter content. Set up Amazon CloudWatch alarms for notification of policy violations.: Guardrails for Amazon Bedrock are designed to enforce compliance and mitigate the risk of inappropriate or sensitive data being generated by foundation models. Guardrails can filter content to prevent sensitive information from being included in model outputs. Amazon CloudWatch alarms can then be configured to notify the company of potential policy violations, making it a comprehensive solution for ensuring privacy and compliance.

upvoted 7 times

psyll0n 4 days, 3 hours ago

C is the correct answer indeed. Supporting documentation.

https://docs.aws.amazon.com/bedrock/latest/userguide/guardrails-sensitive-filters.html

upvoted 1 times

durand26 Most Recent 9 months, 2 weeks ago

Selected Answer: A

A. Use Amazon Macie to scan the model's output for sensitive data and set up alerts for potential violations.

Amazon Macie is an AI service that detects breaches of private information (PII).

On the other hand, Guardrails for Amazon Bedrock is a service used to guard against hallucinations in generative AI apps.

Amazon Macie is much more appropriate if you want to filter out private content.

upvoted 4 times

Rcosmos 1 year ago

Selected Answer: C

Explicação:Guardrails for Amazon Bedrock permite configurar diretrizes que ajudam a impedir que os modelos fundacionais (FMs) incluam informações pessoais ou violem políticas de uso. Isso é feito por meio de filtros e regras de moderação de conteúdo.Amazon CloudWatch pode ser usado para configurar alarmes e notificações quando essas políticas forem violadas, fornecendo a visibilidade necessária em tempo real.

upvoted 1 times

Moon 1 year, 3 months ago

Selected Answer: C

C: Use Guardrails for Amazon Bedrock to filter content. Set up Amazon CloudWatch alarms for notification of policy violations.

Explanation:
Guardrails for Amazon Bedrock allow you to enforce policies that filter out sensitive or inappropriate content, such as personal patient

information, from the model's responses. By configuring these guardrails, you ensure that the model adheres to privacy policies. Additionally, you can set up Amazon CloudWatch alarms to receive notifications when policy violations occur, providing real-time monitoring and alerting.

upvoted 1 times

Gianiluca 1 year, 3 months ago

Selected Answer: C

C. Use Guardrails for Amazon Bedrock to filter content. Set up Amazon CloudWatch alarms for notification of policy violations.

Reasoning:
Requirement:

The company wants to ensure that the disease detection model does not include personal patient information in its responses. This requires a mechanism to filter sensitive information from the model's outputs. They also need a notification system for policy violations.
Solution:

Guardrails for Amazon Bedrock provide built-in content moderation and filtering mechanisms to enforce compliance with policies, such as removing personal information from model outputs.
Amazon CloudWatch Alarms can be used to monitor events or logs (such as detected policy violations) and send notifications when violations occur.

upvoted 1 times

fed6485 1 year, 4 months ago

if the answer A was slightly different..

A. Use Amazon Macie to scan the model's DATA for sensitive data and set up alerts for potential violations.
instead of
A. Use Amazon Macie to scan the model's OUTPUT for sensitive data and set up alerts for potential violations.

as Macie cannot scan the output.. but the data in S3, so A is incorrect.. so C, even if, personal information should always be removed, no point of train on it.

upvoted 4 times

jove 1 year, 4 months ago

Selected Answer: C

Guardrails to prevent, CloudWatch to notify

upvoted 3 times