Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 76 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 76
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

A company is using few-shot prompting on a base model that is hosted on Amazon Bedrock. The model currently uses 10 examples in the prompt. The model is invoked once daily and is performing well. The company wants to lower the monthly cost.
Which solution will meet these requirements?

A. Customize the model by using fine-tuning.

B. Decrease the number of tokens in the prompt.

C. Increase the number of tokens in the prompt.

D. Use Provisioned Throughput.

Show Suggested Answer

---

by Blair77 at *Nov. 12, 2024, 3:45 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

Blair77 Highly Voted 1 year, 4 months ago

Selected Answer: B

Bedrock pricing is based on the number of tokens processed, which includes both input tokens (from the prompt) and output tokens (generated by the model). By decreasing the number of tokens in the prompt, you directly reduce the cost associated with each invocation of the model.

upvoted 6 times

Jessiii Most Recent 1 year, 1 month ago

Selected Answer: B

B. Decrease the number of tokens in the prompt: In a few-shot learning scenario, the number of tokens used in the prompt contributes directly to the cost, as you're billed based on the number of tokens processed during each invocation. By decreasing the number of tokens in the prompt, the company can reduce the cost per invocation while still maintaining the model's performance. This can be done by reducing the number of examples or making the examples more concise.

upvoted 3 times

AzureDP900 1 year, 2 months ago

Selected Answer: D

D. Use Provisioned Throughput

To lower the monthly cost, the company can use Provisioned Throughput (PT) to scale their model's resource utilization. This allows them to pay only for the actual compute time used by the model, rather than paying a fixed monthly fee.

upvoted 1 times

djeong95 1 year, 1 month ago

you are right to point this out but B is a more correct answer. There is a limit as to how much you can save with Provisioned Throughput (given that you were using On Demand before and that the company is okay with a longer term commitment). However, Decrease the number of tokens in the prompt is going to be more effective and doesn't require a longer term commitment.

upvoted 1 times