Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 78 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01
Question #: 78
Topic #: 1
[All AWS Certified AI Practitioner AIF-C01 Questions]

An AI practitioner trained a custom model on Amazon Bedrock by using a training dataset that contains confidential data. The AI practitioner wants to ensure that the custom model does not generate inference responses based on confidential data.
How should the AI practitioner prevent responses based on confidential data?

A. Delete the custom model. Remove the confidential data from the training dataset. Retrain the custom model.

B. Mask the confidential data in the inference responses by using dynamic data masking.

C. Encrypt the confidential data in the inference responses by using Amazon SageMaker.

D. Encrypt the confidential data in the custom model by using AWS Key Management Service (AWS KMS).

Show Suggested Answer

---

by Amitst at *Dec. 5, 2024, 8:58 a.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

Jessiii Highly Voted 1 year, 1 month ago

Selected Answer: A

To ensure that the custom model does not generate inference responses based on confidential data, the best approach is to:

Delete the custom model: If the confidential data was used in training, there's a possibility that the model has memorized this data and might generate it in responses. Removing the model is the first step.
Remove the confidential data from the training dataset: This ensures that confidential information is not included in the model's learning process, mitigating the risk of leakage.
Retrain the custom model: After removing the confidential data, retraining the model with a cleaned dataset ensures that the model does not inadvertently include any sensitive information in its responses.

upvoted 6 times

CollinsWolf Most Recent 7 months, 1 week ago

Selected Answer: B

Dynamic data masking in Amazon Bedrock can be implemented to protect confidential data within inference responses, particularly when using features like Knowledge Bases and Guardrails for RAG (Retrieval Augmented Generation) applications.

upvoted 1 times

65703c1 8 months ago

Selected Answer: A

A is the correct answer

upvoted 1 times

aws4gzone9 11 months, 1 week ago

Selected Answer: A

data memorize is possible.

upvoted 1 times

tcl08 12 months ago

Selected Answer: A

If the model was trained with confidential data, there's a risk it might have memorized that information and could generate it in responses. Deleting the model is the first step to prevent this.

upvoted 2 times

kopper2019 1 year, 1 month ago

A
A. Delete the custom model. Remove confidential data from dataset. Retrain the model.

This is the correct answer because:

Once a model learns from confidential data, that information becomes embedded in its parameters
The only way to truly prevent it from using that knowledge is to retrain from scratch without the confidential data
Deleting and retraining ensures the model has no access to the sensitive information

upvoted 2 times

Moon 1 year, 3 months ago

Selected Answer: A

A: Delete the custom model. Remove the confidential data from the training dataset. Retrain the custom model.

Explanation:
If the training dataset contains confidential data, the model may inadvertently learn and generate responses based on that data. The only way to ensure that the model does not generate responses based on the confidential data is to:

Remove the confidential data from the training dataset.
Retrain the custom model using the updated dataset.
This process ensures that the model is not influenced by the sensitive information.

upvoted 4 times

BhaskarSadineni 1 year, 3 months ago

Selected Answer: A

Explanation:
Once a model is trained, the data used for training is embedded in its parameters. If confidential data is included in the training dataset, it can influence the responses the model generates.
Simply masking or encrypting inference responses will not ensure the model doesn't generate responses derived from the confidential data; the issue originates in the training process itself.

upvoted 3 times

may2021_r 1 year, 3 months ago

Selected Answer: A

The correct answer is A. Once a model is trained on confidential data, it must be retrained without it.

upvoted 2 times

AKG85 1 year, 3 months ago

Selected Answer: A

Delete the custom model, remove the confidential data, and retrain the model is the best approach because it ensures that the model will not retain or generate responses based on any confidential information

upvoted 3 times

ap6491 1 year, 3 months ago

Selected Answer: A

Once a model is trained on data, its outputs may inherently reflect patterns or details derived from the training dataset, including confidential data.
To ensure the custom model does not generate inference responses based on confidential data, the only reliable solution is to:
- Remove the confidential data from the training dataset.
- Retrain the model with the updated dataset.

This approach ensures the model is not influenced by sensitive information during inference.

Option B is incorrect.

Dynamic data masking hides sensitive information in database query results or outputs but does not prevent the model from generating responses influenced by the confidential data. The model would still "know" the sensitive patterns.

upvoted 2 times

Dandelion2025 1 year, 3 months ago

Selected Answer: B

The company should mask the confidential information

upvoted 2 times

Amitst 1 year, 4 months ago

Selected Answer: B

This is the most efficient method, effectively maintaining data privacy and security.

upvoted 2 times