Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 98 DIS...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 98

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

A company wants to implement a large language model (LLM) based chatbot to provide customer service agents with real-time contextual responses to customers' inquiries. The company will use the company's policies as the knowledge base.

Which solution will meet these requirements MOST cost-effectively?

A. Retrain the LLM on the company policy data.

B. Fine-tune the LLM on the company policy data.

C. Implement Retrieval Augmented Generation (RAG) for in-context responses.

D. Use pre-training and data augmentation on the company policy data.

Show Suggested Answer

---

by aws_Tamilan at *Dec. 28, 2024, 4:56 a.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

65703c1 8 months ago

Selected Answer: C

C is the correct answer

upvoted 1 times

Jessiii 1 year, 1 month ago

Selected Answer: C

Retrieval Augmented Generation (RAG) is a technique that combines the power of pre-trained large language models (LLMs) with a retrieval mechanism to fetch relevant documents or data (like the company's policies) and incorporate them into the response generation process. This approach allows the model to produce accurate, context-aware responses without the need to retrain or fine-tune the model extensively on the entire policy dataset, making it a cost-effective and efficient solution.

upvoted 2 times

85b5b55 1 year, 2 months ago

Selected Answer: C

RAG provides most cost-effective solution

upvoted 1 times

may2021_r 1 year, 3 months ago

Selected Answer: C

The correct answer is C. RAG allows direct use of policy documents without expensive model training.

upvoted 1 times

aws_Tamilan 1 year, 3 months ago

Selected Answer: C

RAG (Option C) is the most cost-effective choice because it allows the LLM to dynamically retrieve relevant information from a predefined knowledge base (the company policy) at inference time, without needing extensive fine-tuning or retraining of the model. This reduces the need for costly computational resources while still providing accurate, contextual responses.

upvoted 1 times