Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 119 DI...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 119

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

A company has developed an ML model to predict real estate sale prices. The company wants to deploy the model to make predictions without managing servers or infrastructure.

Which solution meets these requirements?

A. Deploy the model on an Amazon EC2 instance.

B. Deploy the model on an Amazon Elastic Kubernetes Service (Amazon EKS) cluster.

C. Deploy the model by using Amazon CloudFront with an Amazon S3 integration.

D. Deploy the model by using an Amazon SageMaker endpoint.

Show Suggested Answer

---

by aws_Tamilan at *Dec. 28, 2024, 5:55 a.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

65703c1 8 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

Jessiii 1 year, 1 month ago

Selected Answer: D

Fully managed service for deploying ML models, allowing predictions without managing infrastructure.

upvoted 2 times

may2021_r 1 year, 3 months ago

Selected Answer: D

The correct answer is D. Deploying the model using an Amazon SageMaker endpoint allows for serverless predictions.

upvoted 1 times

aws_Tamilan 1 year, 3 months ago

Selected Answer: D

D. Deploy the model by using an Amazon SageMaker endpoint.

Explanation:

Amazon SageMaker is a fully managed service that enables you to quickly build, train, and deploy machine learning models at scale. Deploying a model using an Amazon SageMaker endpoint allows the company to make predictions without needing to manage servers or infrastructure. SageMaker automatically handles the provisioning of resources, scaling, and maintenance, making it an ideal solution for production-grade ML deployments.

upvoted 2 times