Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 125 DI...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 125

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

HOTSPOT

-

A company has developed a large language model (LLM) and wants to make the LLM available to multiple internal teams. The company needs to select the appropriate inference mode for each team.

Select the correct inference mode from the following list for each use case. Each inference mode should be selected one or more times.

Answer Area

The company's chatbot needs predictions from the LLM to understand users' intent with minimal latency.	Select... ▼ (Select... / Batch transform / Real-time inference)
A data processing job needs to query the LLM to process gigabytes of text files on weekends.	Select... ▼ (Select... / Batch transform / Real-time inference)
The company's engineering team needs to create an API that can process small pieces of text content and provide low-latency predictions.	Select... ▼ (Select... / Batch transform / Real-time inference)

Show Suggested Answer

by LonghornFan at *Feb. 10, 2025, 11:09 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

65703c1 8 months ago

Below is the correct answer:

Realtime
Batch
Realtime

upvoted 1 times

4729e6c 1 year, 1 month ago

In Scenario 1, the company's chatbot requires immediate predictions to accurately interpret user intent during interactions. Implementing real-time inference ensures that the model processes each user input instantaneously, providing timely and relevant responses essential for effective communication.

For Scenario 2, the data processing job involves handling large volumes of text data during specific periods, such as weekends. Utilizing batch transform allows the system to process this data in bulk, optimizing resource utilization and efficiency without the need for immediate results.

In Scenario 3, the engineering team aims to develop an API capable of swiftly processing small text inputs and delivering prompt predictions. Adopting real-time inference enables the API to handle each request as it arrives, ensuring low-latency responses critical for user-facing applications.

upvoted 1 times

Jessiii 1 year, 1 month ago

real time
batch
real time
chatbot (real-time, low-latency predictions): Real-time inference
Data processing job (large datasets on weekends): Batch transform

upvoted 1 times

LonghornFan 1 year, 1 month ago

Scenario 1: The company’s chatbot needs predictions from the LLM to understand users’ intent with minimal latency. Answer: Real-time inference
Scenario 2: A data processing job needs to query the LLM to process gigabytes of text files on weekends. Answer: Batch transform
Scenario 3: The company’s engineering team needs to create an API that can process small pieces of text content and provide low-latency predictions. Answer: Real-time inference

upvoted 2 times