Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 167 DI...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 167

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

A publishing company built a Retrieval Augmented Generation (RAG) based solution to give its users the ability to interact with published content. New content is published daily. The company wants to provide a near real-time experience to users.

Which steps in the RAG pipeline should the company implement by using offline batch processing to meet these requirements? (Choose two.)

A. Generation of content embeddings

B. Generation of embeddings for user queries

C. Creation of the search index

D. Retrieval of relevant content

E. Response generation for the user

Show Suggested Answer

---

by TeeMal at *June 6, 2025, 12:32 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

pgcity3 3 weeks, 3 days ago

Selected Answer: C

A and C

upvoted 1 times

d7797d1 2 months, 1 week ago

Selected Answer: C

C is also the answer

upvoted 1 times

rstrstrst 6 months, 1 week ago

Selected Answer: A

The correct options are A. Generation of content embeddings and C. Creation of the search index.

upvoted 1 times

AZD98 7 months, 2 weeks ago

Selected Answer: C

Its both A and C

upvoted 2 times

Dubez 7 months, 2 weeks ago

Selected Answer: A

The correct answers are A. Generation of content embeddings and C. Creation of the search index.

Here's why these are the best choices:

A. Generation of content embeddings:

Can be done as a batch process when new content is published
Computationally intensive task
Not time-critical for user interaction
Can be scheduled periodically
Prepares content for efficient retrieval
C. Creation of the search index:

Can be built offline using the generated embeddings
Resource-intensive process
Can be updated periodically
Crucial for efficient retrieval
Doesn't need to happen in real-time

upvoted 2 times

65703c1 8 months ago

Selected Answer: A

AC is the correct answer

upvoted 1 times

rinip86277 9 months ago

Selected Answer: A

A and C is correct I think

upvoted 1 times

neil1985_jy 9 months, 1 week ago

Selected Answer: A

A and C
1. Documents are published daily, the company can periodically process and embed this content in batches
2. Once embeddings are generated, they need to be indexed in a vector database

upvoted 1 times

Rcosmos 9 months, 3 weeks ago

Selected Answer: U

✅ Etapas que devem ser feitas offline (em lote):
A. Geração de incorporações de conteúdo
Isso é feito quando novos conteúdos são adicionados. Como os documentos não mudam frequentemente depois de publicados, essa etapa pode ser feita em lote.

C. Criação do índice de pesquisa
Após gerar as incorporações dos documentos, é necessário indexá-los para permitir busca eficiente. Isso também pode ser feito em lote, e atualizado conforme novos conteúdos são publicados.

upvoted 1 times

026dda3 10 months ago

Selected Answer: A

A C * A. Generation of content embeddings: Creating embeddings for all the published content is a computationally intensive process that doesn't need to happen in real-time as users interact with the system.
* C. Creation of the search index: The search index, typically a vector database in a RAG system, needs to be built and updated to store the content embeddings.

upvoted 3 times

TeeMal 10 months ago

Selected Answer: A

Correct:
Content Embedding (A) – Transform published documents into embeddings using a model (e.g., via Amazon Titan or OpenSearch ML). This can be done offline in batches, especially for static or periodically updated content like daily publications.

Search Index Creation (C) – After generating embeddings, these need to be indexed (e.g., using Amazon OpenSearch or FAISS). This step can also be handled offline, as it's only needed when content updates.

Wrong:
B, D, and E require real-time processing for a near-real-time experience.

upvoted 2 times