Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 172 DI...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 172

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

What is the purpose of chunking in Retrieval Augmented Generation (RAG)?

A. To avoid database storage limitations for large text documents by storing parts or chunks of the text

B. To improve efficiency by avoiding the need to convert large text into vector embeddings

C. To improve the contextual relevancy of results retrieved from the vector index

D. To decrease the cost of storage by storing parts or chunks of the text

Show Suggested Answer

---

by Zangi at *July 16, 2025, 3:25 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

psyll0n 3 days ago

Selected Answer: C

C. Correct answer and the only one that makes sense in this context.

upvoted 1 times

SkillsFuture 3 months, 4 weeks ago

Selected Answer: C

C. To improve the contextual relevancy of results retrieved from the vector index
✅ Correct. Chunking ensures that queries retrieve smaller, semantically relevant portions of text instead of entire documents, improving accuracy and relevance of responses.

upvoted 1 times

samplunk 4 months, 1 week ago

Selected Answer: C

In Retrieval Augmented Generation (RAG), chunking splits large documents into smaller, meaningful pieces before generating embeddings. This is done to:
- Ensure that retrieved chunks are contextually relevant to the user query
- Improve the accuracy and quality of responses from the generative model
- Avoid situations where a single embedding for a large document loses fine-grained context
Other options are incorrect:
A. Database storage limitations → Chunking is not primarily for storage.
B. Avoiding embeddings → Chunking actually requires creating embeddings for each chunk.
D. Decrease cost of storage → Not the main purpose; it's about contextual relevancy.

upvoted 1 times

Liongeek 7 months ago

Selected Answer: B

The purpose of chunking in Retrieval-Augmented Generation (RAG) is to break large documents into smaller, manageable pieces to improve the efficiency, accuracy, and scalability of the information retrieval process

upvoted 2 times

hype23 8 months, 3 weeks ago

Selected Answer: C

Retrieval Augmented Generation (RAG) systems enhance Large Language Model (LLM) responses by providing relevant external knowledge. A fundamental step in building effective RAG systems is chunking, the process of dividing large documents into smaller, digestible pieces.

upvoted 4 times