Exam AWS Certified AI Practitioner AIF-C01 All Questions

View all questions & answers for the AWS Certified AI Practitioner AIF-C01 exam

Go to Exam

EXAM AWS CERTIFIED AI PRACTITIONER AIF-C01 TOPIC 1 QUESTION 217 DI...

Exam question from Amazon's AWS Certified AI Practitioner AIF-C01

Question #: 217

Topic #: 1

[All AWS Certified AI Practitioner AIF-C01 Questions]

A bank is fine-tuning a large language model (LLM) on Amazon Bedrock to assist customers with questions about their loans. The bank wants to ensure that the model does not reveal any private customer data.

Which solution meets these requirements?

A. Use Amazon Bedrock Guardrails.

B. Remove personally identifiable information (PII) from the customer data before fine-tuning the LLM.

C. Increase the Top-K parameter of the LLM.

D. Store customer data in Amazon S3. Encrypt the data before fine-tuning the LLM.

Show Suggested Answer

---

by nand2804 at *Aug. 4, 2025, 5:30 p.m.*

Disclaimers:
- ExamTopics website is **not** related to, affiliated with, endorsed or authorized by Amazon.
- Trademarks, certification & product names are used for reference only and belong to Amazon.

Comments

nand2804 Highly Voted 8 months ago

Selected Answer: B

When fine-tuning a large language model (LLM) with customer data, it is essential to ensure data privacy and compliance with regulations (like GDPR or HIPAA). The most effective and direct solution to prevent the model from learning or exposing sensitive customer information is to:

✅ Remove personally identifiable information (PII) from the dataset before fine-tuning.

This helps:

Prevent the model from memorizing or leaking private data

Reduce privacy and compliance risks

Follow best practices for data minimization

upvoted 5 times

chris_spencer Most Recent 2 months ago

Selected Answer: B

remove Pii

upvoted 1 times

Sesh_14 2 months, 3 weeks ago

Selected Answer: B

Explanation:

The most reliable way to ensure that an LLM does not reveal private customer data is to prevent that data from ever being learned by the model.

Why B is correct:

If PII is included in fine-tuning data, the model can memorize and later reproduce it

Removing or anonymizing PII before fine-tuning is a best practice for privacy, compliance, and risk reduction

This approach is preventive, not reactive, and works regardless of how the model is later prompted

upvoted 1 times

ticoY_2025 3 months, 2 weeks ago

Selected Answer: A

Amazon Bedrock Guardrails are specifically designed to prevent sensitive or disallowed content from being generated or returned by an LLM, including private customer data and PII.

upvoted 1 times

Hacksawz 4 months, 1 week ago

Selected Answer: B

The correct answer is:

B. Remove personally identifiable information (PII) from the customer

data before fine-tuning the LLM.

Removing PII before fine-tuning ensures that the model never learns or has access to sensitive customer data, fully protecting privacy and compliance. Guardrails (A) help filter outputs but cannot prevent private data from being encoded in the model during fine-tuning. Adjusting model parameters (C) and encrypting data in S3 (D) protect data at rest or impact inference, but neither addresses the foundational risk: training on sensitive information itself.

upvoted 2 times

samplunk 4 months, 1 week ago

Selected Answer: B

When fine-tuning an LLM, any information in the training data can be memorized by the model.
To prevent the model from revealing private customer information, the most important step is to sanitize the dataset by removing PII before fine-tuning.

A. Amazon Bedrock Guardrails – helps control model outputs after deployment but does not prevent the model from memorizing PII during fine-tuning. Works on outputs, not training data

upvoted 1 times

rstrstrst 6 months, 1 week ago

Selected Answer: A

The answer is A, the model still needs access to the customer data to be able to answer questions. You cannot remove it.

upvoted 3 times

allasiki 7 months ago

Selected Answer: A

A. Use Amazon Bedrock Guardrails.
Guardrails in Bedrock let you enforce policies on what the model can output, including blocking personally identifiable information (PII) or other sensitive content. This directly addresses the requirement to ensure the model does not reveal private customer data.

B. Remove personally identifiable information (PII) from the customer data before fine-tuning the LLM.
This is a good data hygiene practice, but the question specifically emphasizes ensuring the model does not reveal private data. Even if PII is removed, the model might still generate unsafe or unintended outputs without guardrails.

upvoted 2 times

7ebc5c4 7 months, 1 week ago

Selected Answer: A

I believe guardrails are better placed to do that

upvoted 2 times